



# AaaS — Demo Script

The live demo is the single most important thing judges remember. This script exists so the presenter can do it under pressure, with background noise, with a dodgy projector, and with one hand holding a water bottle.

Everything in this script is **shipped and rehearsable today** — every scene is a real feature running on this laptop. There is no vapourware scene. If a scene fails twice in rehearsal, cut the scene, not the honesty.

## Setup before judges arrive

1. Laptop plugged in, 100% battery, notifications silenced.
2. Start the stack: double-click `START-DEMO.bat` in `C:\SUBARNAREKHA` (from the repo: `start-dev.bat`). Wait for "DEMO IS READY" — about 1 minute; gateway :8000, TTS :8001, STT :8002, translate :8003.
3. **Pre-warm the caches:** open `http://127.0.0.1:8000/demo/ssepd-odisha/` and run one Read-aloud, one Translate, and one Read-document pass. First-time synthesis takes seconds; cached replays are instant. Judges get the cached experience. Translations persist on disk (`services/translate/cache/translations.sqlite3`), so pages warmed in rehearsal stay instant across restarts — a cold whole-page translate of SSEPD takes ~2 min on CPU, a warm one 0.2 s. **Never delete that file the night before.** If you edit a demo page's text, re-run Translate on it once to re-warm.
4. Browser tabs, in order: - **Tab 1:** `http://127.0.0.1:8000/demo/ssepd-odisha/` — the flagship: SSEPD department (pensions, UDID, disability schemes). Scenes 1–4 and 6 all run here — scanned notice + 2-page PDF + application form on one page. - **Tab 2:** `http://127.0.0.1:8000/demo/bse-odisha/` — spare: BSE board mirror with the exam circular, in case a judge asks for a second site. - **Tab 3:** `https://ssepd.odisha.gov.in/` — the REAL SSEPD department site, with the browser extension loaded (`chrome://extensions` → verify it's on). **Not Jajpur for voice:** since 2026 jajpur.odisha.gov.in sends `Permissions-Policy: microphone=()`, which blocks every mic on its pages. Jajpur still works for Read aloud / Translate / Read document. Extension settings for the booth (verified 2026-09-28): popup → **On-device mode OFF** (uses the full local models); first visit to a real site Chrome asks twice — microphone and "access other apps/devices on this device" (local network) — click **Allow** both. Do one Read aloud + Translate on Tab 3 during setup so they're cached.
5. External microphone plugged in and tested — the built-in mic struggles with booth noise. One headphone bud in your ear for private TTS checks.
6. Insurance: a screen recording of a full successful run, in a hidden tab.

## The one-liner (say it before touching anything)

"Every Odisha government site is written in English, for sighted people who type. Four crore Odias — the elderly, the blind, those who can't read English — live on the other side of that wall. AaaS is one button that tears the wall down — on any website, in Odia, by voice. Watch."

## Scene 1 — Read aloud (60 seconds)

*Tab 1, the SSEPD site. Click the floating ଥି button. The icon-grid panel opens: six tiles, Odia labels first.*

"This is the disability-welfare department's portal — pensions, UDID cards. Its whole audience is people the web serves worst. Meet a sixty-eight-year-old pensioner who can't read small English text."

*Tap ଠିକ୍ ଶୁଣାଅି · Read aloud. The Odia voice starts; each block gets an amber highlight as it's spoken.*

"The page reads itself aloud — and highlights exactly what it's reading, so a low-vision or dyslexic reader can follow along. These ⏮ ⏭ buttons skip or replay a section. Space pauses. Escape stops."

*Skip forward once with ⏭ so judges see the highlight jump. Press Esc.*

"The voice is Meta's MMS Odia model running on this laptop — no cloud. We taught it to say numbers: '୩୦୦ ଟଙ୍କା' comes out as 'three hundred rupees' in Odia words, because the raw model goes silent on digits."

## Scene 2 — Translate + Easy Read (60 seconds)

*Tap ଅନୁବାଦ · Translate → Odia. The page rewrites itself in Odia, in place — layout intact.*

"Whole-page translation to Odia — AI4Bharat's IndicTrans2, running locally. Notice the ↶ undo button on the tile: one tap restores the original. Nothing is destructive."

*Tap ସରଳ ଠିକ୍ · Easy Read.*

"Easy Read simplifies the bureaucratic language — 'shall furnish the requisite documents' becomes 'must give the papers'. It's rule-based, fully offline, and it simplifies in whatever language the page is currently in — here, Odia."

## Scene 3 — Read document: the scanned circular (90 seconds — the peak)

*Scroll to "Notices & Scanned Documents". Point at the scanned pension notification and the "UDID enrolment camp circular (scanned PDF, 2 pages)".*

"Now the hard problem. Government notices are scans — photographs of paper. Screen readers see nothing. Translators see nothing. For a blind citizen, this circular does not exist."

*Tap ଚିତ୍ର ଠିକ୍ · Read document. The chooser asks which document — it found both the scanned image and the PDF; tap the 2-page UDID circular.*

"It found the scanned PDF, ran OCR on it — Tesseract, on this laptop — simplified it, translated it to Odia..."

*The side-by-side view opens: original scanned pages LEFT, Odia text RIGHT. It starts reading aloud; the page being spoken glows amber on both sides.*

"...and now it reads it aloud in Odia, showing you the original page and the Odia text side by side. Same stamp, same reference number, both languages. The amber highlight follows the voice across pages."

*Let it speak one page. Tap  to maximize briefly. Close.*

## Scene 4 — Speak to fill: the pension application (90 seconds)

*Same tab. Scroll to "Apply for Madhu Babu Pension".*

"Arun is sixty-eight, arthritic hands, speaks only Odia. His pension just went up — he needs to apply. Watch him fill this form without touching the keyboard."

*Tap କହି ଲେଖି. Speak to fill with no field selected. The guided mode starts: it highlights the name field, ASKS for it aloud in Odia, beeps, and listens.*

*Speak, one field per prompt:*

- Name: "ପୂର୍ଣ୍ଣଚନ୍ଦ୍ର ମହାନ୍ତି" → lands as "Purnnachandra Mahanti"
- Village: "ଗ୍ରାମ ଧର୍ମଶାଳା" → "Grama Dharmashala"
- Phone: "ନଅ ଆଠ ସାତ ଛଅ ଦୁଇ ପାଞ୍ଚ..." (spoken digits, "double five") → 987655...
- Age: "ଷାଠିଏ ପାଞ୍ଚ" (sixty-five as words) → 65
- Email: "purna at gmail dot com" → purna@gmail.com
- Details: one Odia sentence about the disability → romanized text

*(Staying silent on a field just skips it — rehearse the full six so the counter reads 6/6.)*

"Three things just happened that are genuinely hard. One: his name was **transliterated, not translated** — Purnnachandra means 'full moon', and a naive pipeline writes 'full moon' in the name box. Two: spoken numbers became digits — the speech model writes 'nine eight seven' as words. Three: it never presses Submit. It fills, it shows a ✓ for each answer, and it leaves the final decision to Arun. Progress counter, retry on a misheard number, Escape to stop — this is a kiosk-grade flow."

*Press Enter on the focused Submit button. The demo confirmation appears.*

## Scene 5 — Voice command on the REAL site (60 seconds)

*Tab 3 — the real ssepd.odisha.gov.in, extension loaded.*

"Everything so far ran on mirrors. This is the actual live website of the SSEPD department — we don't control it. The extension injects the same button."

*Tap କହି ଚଳାଅ. Voice command. Speak: "ବିଜ୍ଞପ୍ତି" (notifications). Other verified commands: "ବଜେଟ୍", "ଅତିରିକ୍ତ ଲାଇଭ୍‌ରେଟ୍", "ଯୋଗାଯୋଗ".*

"Speak in Odia — we match what you said against every link on the page, including ones hidden inside menus, and English link names are matched phonetically, the way an Odia speaker says them."

*It highlights and opens Notifications. If it shows the "did you mean" buttons instead, SAY SO — it's a feature:*

"When it isn't sure, it never guesses — it shows the top candidates as buttons. Voice gets you 90% there, one tap finishes it. No dead ends."

## Scene 6 — Comfortable letters (30 seconds, closer)

Back on any page. Open the panel's Reading comfort section, tick ଥିବା ଅକ୍ଷର · Comfortable letters.

"For dyslexic readers: OpenDyslexic letterforms — bottom-weighted so letters don't flip — extra spacing, warm background. Odia script keeps its proper conjunct rendering. And notice: our own panel is unaffected — we hold our product to the standard we're fixing."

Close with the pitch line:

"One widget, one gateway, four on-prem AI services, works on any website — the government doesn't have to rebuild a single portal. Bhashini gives India language AI. AaaS turns it into hands, eyes and ears for the citizens who need it most. Odisha first, India next."

## Timing budget

| SCENE                         | TIME         |
|-------------------------------|--------------|
| One-liner                     | 0:20         |
| 1 Read aloud                  | 1:00         |
| 2 Translate + Easy Read       | 1:00         |
| 3 Scanned circular            | 1:30         |
| 4 Voice form fill             | 1:30         |
| 5 Real site + voice command   | 1:00         |
| 6 Comfortable letters + close | 0:40         |
| <b>Total</b>                  | <b>~7:00</b> |

Short-on-time order of cuts: Scene 6 → Scene 2's Easy Read half → Scene 1's skip-button beat. **Never cut Scene 3 or 4** — they are the demo.

## Fallbacks

- **Mic fails / booth too loud:** every voice scene has a typed equivalent — the widget accepts simulated input events; rehearse the recorded-video fallback instead. Do NOT struggle with a mic twice in front of judges.
- **A service dies:** every engine has a deterministic mock fallback — the widget stays alive and says so in the status pill. Restart via `start-dev.bat` ; services come back in under a minute (TTS needs ~40 s).
- **Internet dies:** nothing in Scenes 1–4 and 6 needs internet. Scene 5 (real site) does — swap to Tab 1's mirror of the same site and say "this mirror is byte-identical to the live site we can't reach from here."

- **Real site says "Mic unavailable":** the site bans microphones (Jajpur does, via its server). Say so — "this site's own security policy blocks microphones; here's the same command on SSEPD" — and switch to Tab 3.
- **Everything dies:** the hidden tab has the full recorded run. Narrate over it with the same script.

## What NOT to claim

- Don't name models we don't run. We run: **IndicWav2Vec** (STT, AI4Bharat), **IndicTrans2** (translation, AI4Bharat), **Meta MMS-TTS** (speech), and **Tesseract** (OCR). No Whisper, no IndicConformer, no Piper, no LLM.
- The sign-language avatar and the exam module are **not features**. If asked: avatar is a roadmap concept (a three.js proof-of-concept exists), the exam module was deliberately cut for focus.
- Don't quote latency numbers from memory — say "first synthesis takes a few seconds on CPU; everything repeated is cached and instant", then show it.